

Julia on Fugaku

Mosè Giordano
m.giordano@ucl.ac.uk



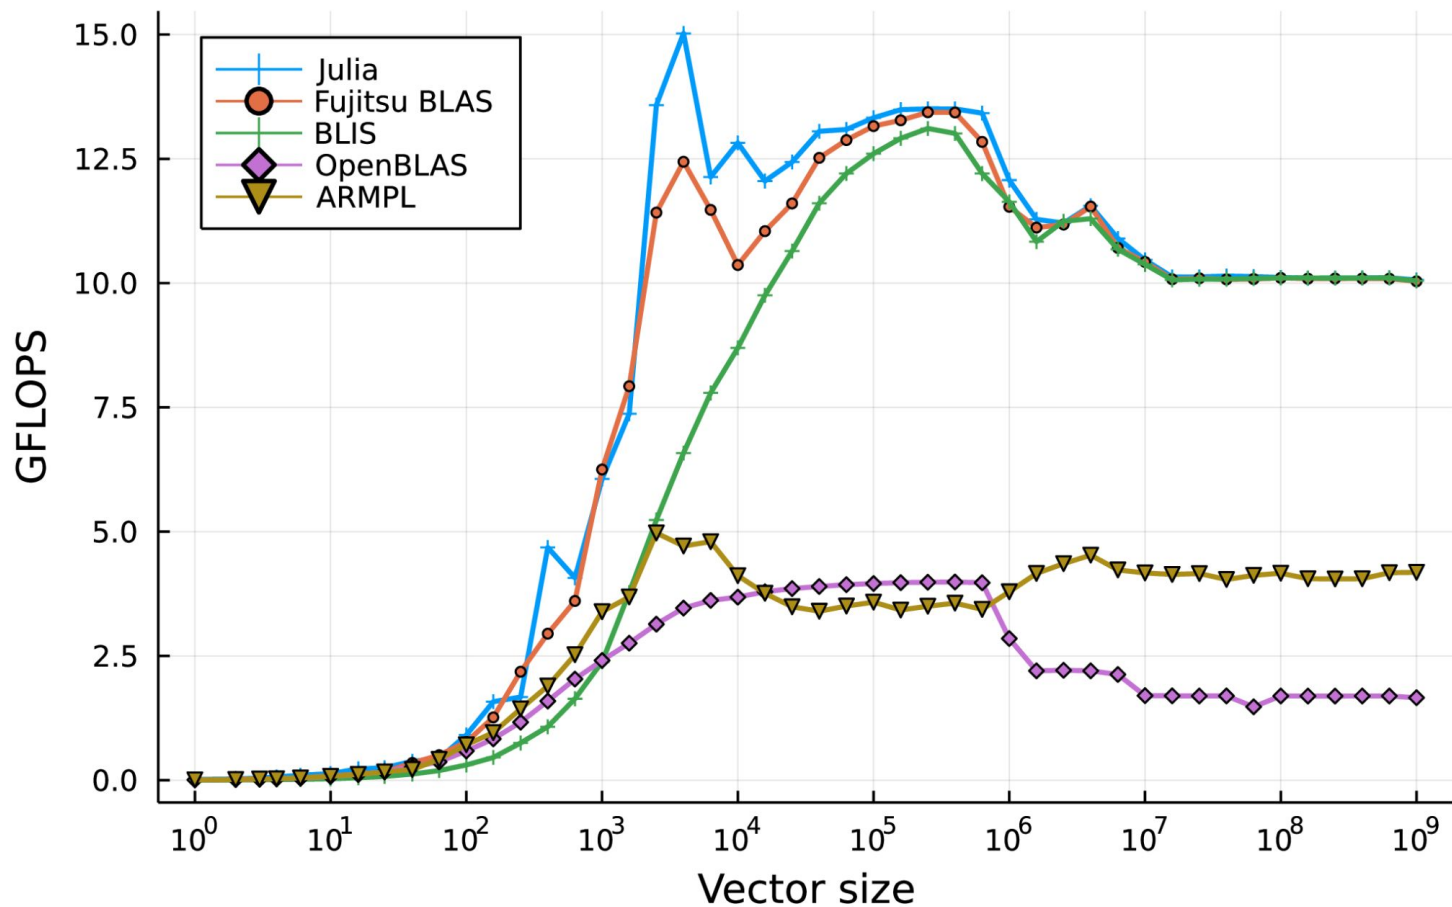
Level 1 BLAS showdown

```
function axpy!(a, x, y)
    @simd for i in eachindex(x, y)
        @inbounds y[i] = muladd(a, x[i], y[i])
    end
    return y
end
```

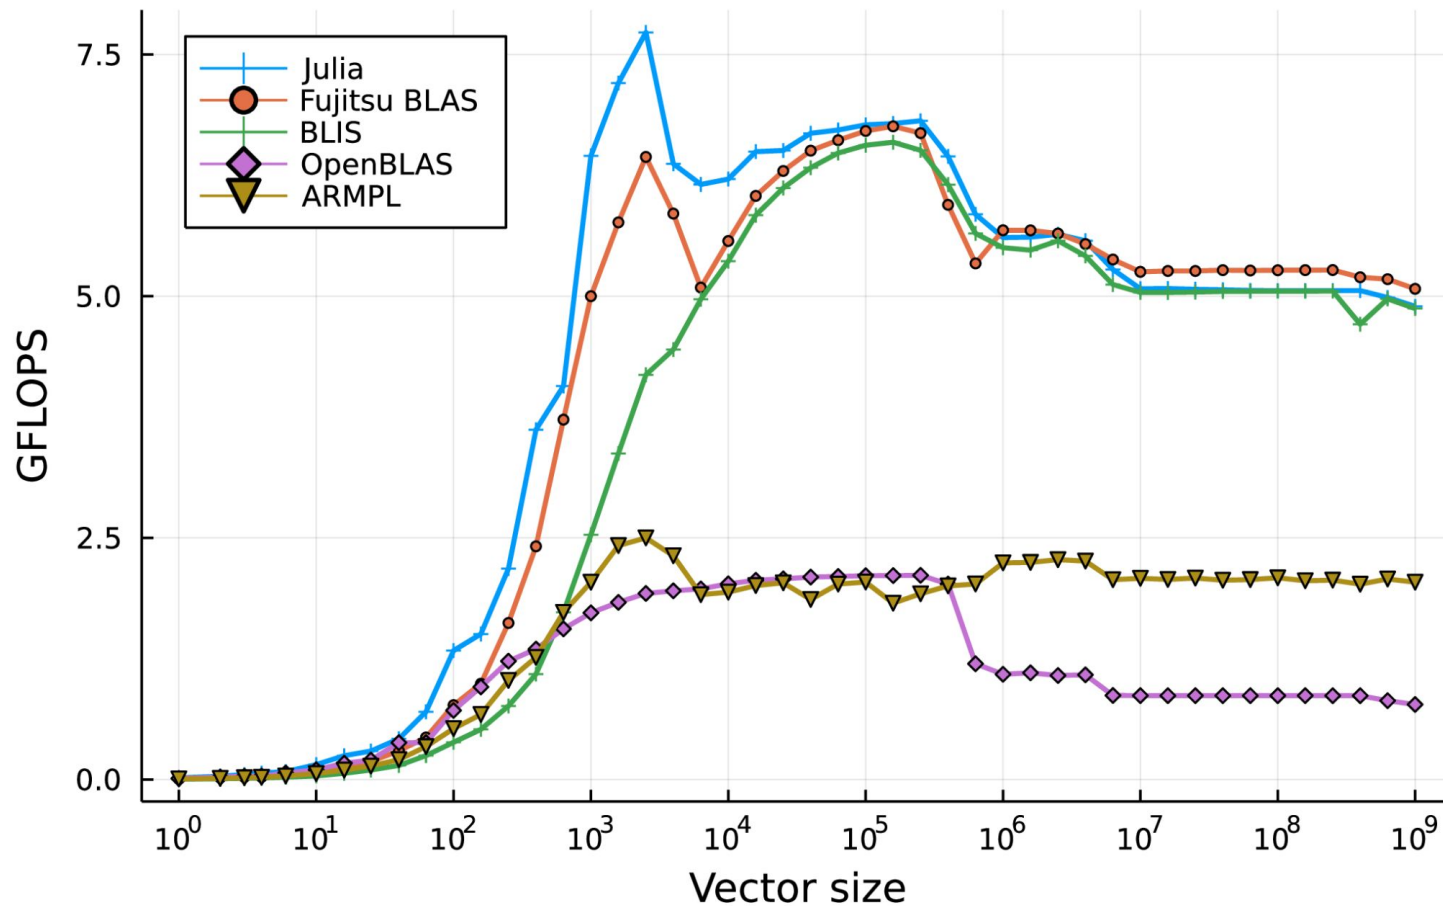
vs

```
LinearAlgebra.BLAS.axpy!(a, x, y)
```

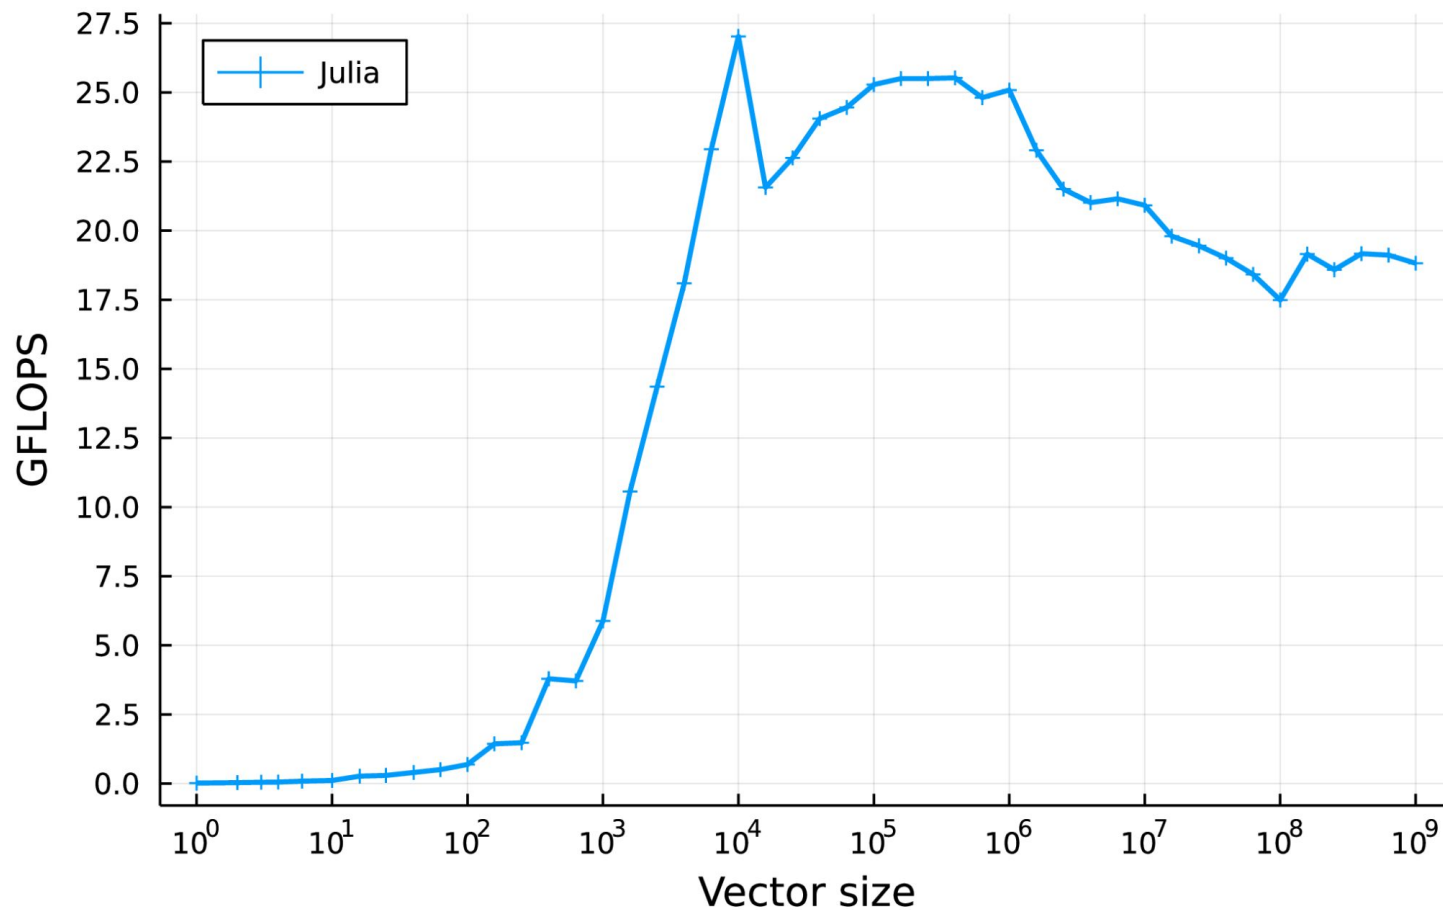
axpy (single precision)



axpy (double precision)



axpy (half precision)




MPI.jl

- Low-level access to MPI
- High-level convenience wrappers
- Deals with MPI ABI


One of the oldest Julia packages (2012)

```
function pingpong(T::Type, bufsize::Int,
                  iters::Int, comm::MPI.Comm)
    rank = MPI.Comm_rank(comm)
    buffer = zeros(T, bufsize)
    tag = 0
    MPI.Barrier(comm)
    tic = MPI.Wtime()
    for i in 1:iters
        if iszero(rank)
            MPI.Send(buffer, comm; dest=1, tag)
            MPI.Recv!(buffer, comm; source=1, tag)
        elseif isone(rank)
            MPI.Recv!(buffer, comm; source=0, tag)
            MPI.Send(buffer, comm; dest=0, tag)
        end
    end
    toc = MPI.Wtime()
    return (toc - tic) / iters
end
```


Latency of MPI PingPong @ Fugaku




Throughput of MPI PingPong @ Fugaku




Latency of MPI Allreduce @ Fugaku (384 nodes, 1536 ranks)



Latency of MPI Gatherv @ Fugaku (384 nodes, 1536 ranks)



Latency of MPI Reduce @ Fugaku (384 nodes, 1536 ranks)



Opportunities for improvements

- Julia is not optimised for A64FX, but every version gets better thanks to upstream improvements in LLVM
 - Keep them coming!
- Compilation latency on A64FX hits particularly hard a JIT language
 - There are ongoing works to continue reducing compilation latency in Julia and to improve static compilation story
- Custom reductions don't work in MPI.jl on non-Intel architectures
 - Can be fixed, but it needs someone to do the work
- ~~Runtime detection of hardware support for Float16~~ Done!
 - Recently tested also on AVX512-FP16